


OutputFormat

Validate output specifications

RecordWriter implementation to write output files of the job



REDUCER

- Reduce function will take Key Value pairs from multiple Map functions as input and Reduce them to output
- Keys are grouped with values. Reduce function is called once per key and its values.
- There could be 0, 1 or more Reduce function for a MapReduce job



InputFormat

Validate inputs

Input files into logical InputSplits

RecordReader implementation to extract logical records




DISSECTING MAPREDUCE PROGRAM

Writable

A serializable object which implements a simple, efficient, serialization protocol

Fast, compact & effective



www.hadoopinrealworld.com

MAPPER

- Dataset is divided in to multiple parts - Input Splits
- Each Mapper process an Input Split
- Each Mapper can be called multiple times depending on the content of Input Split
- Mapper will emit Key Value pairs as output
- There will be one or more Mapper in a MapReduce job




DISSECTING MAPREDUCE PROGRAM

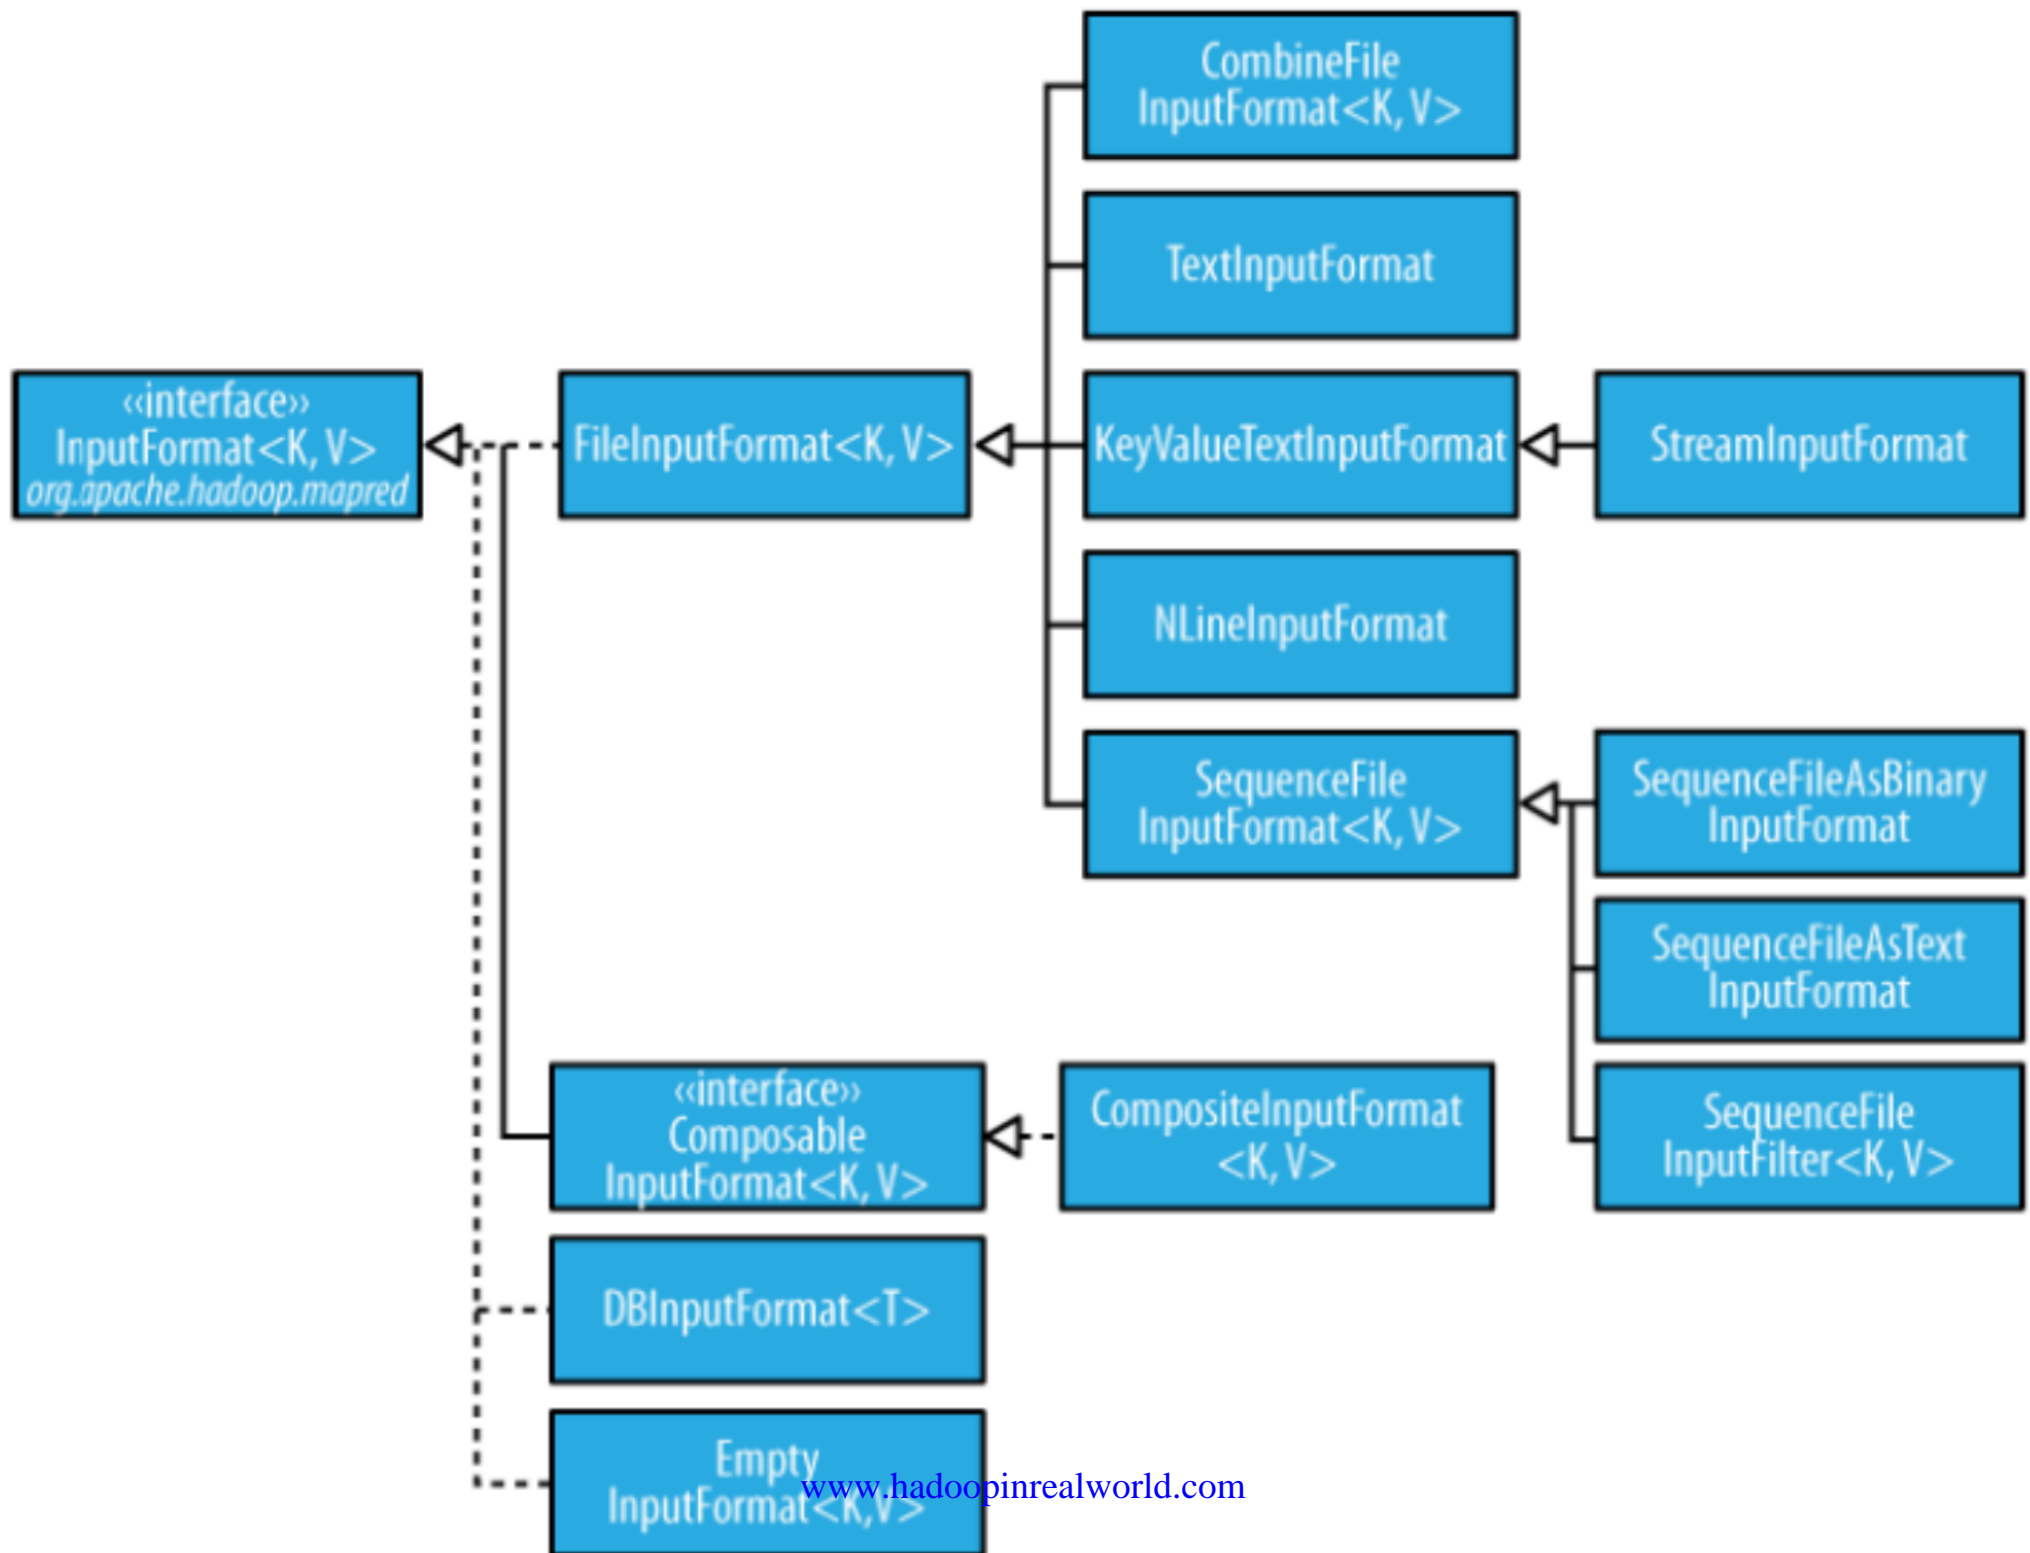
InputFormat

Validate inputs

Input files into logical InputSplits

RecordReader implementation to extract logical records






InputFormat

Validate inputs

Input files into logical InputSplits

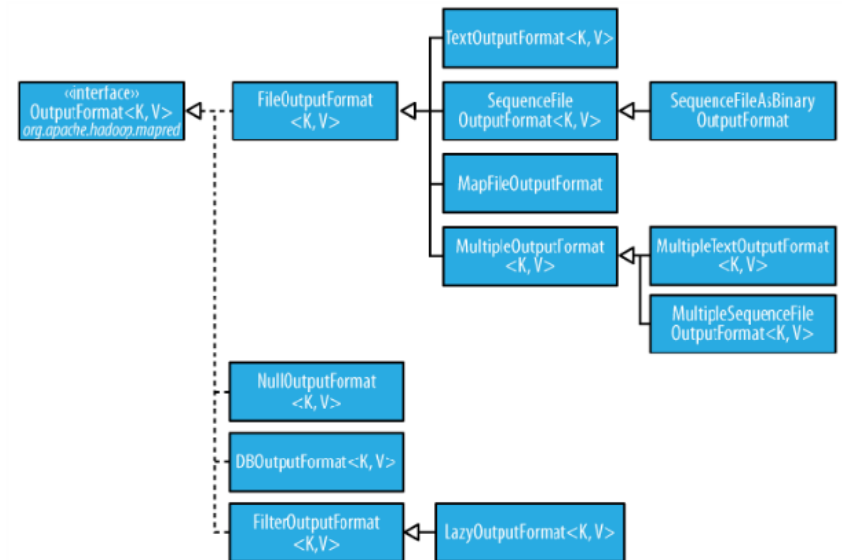
RecordReader implementation to extract logical records

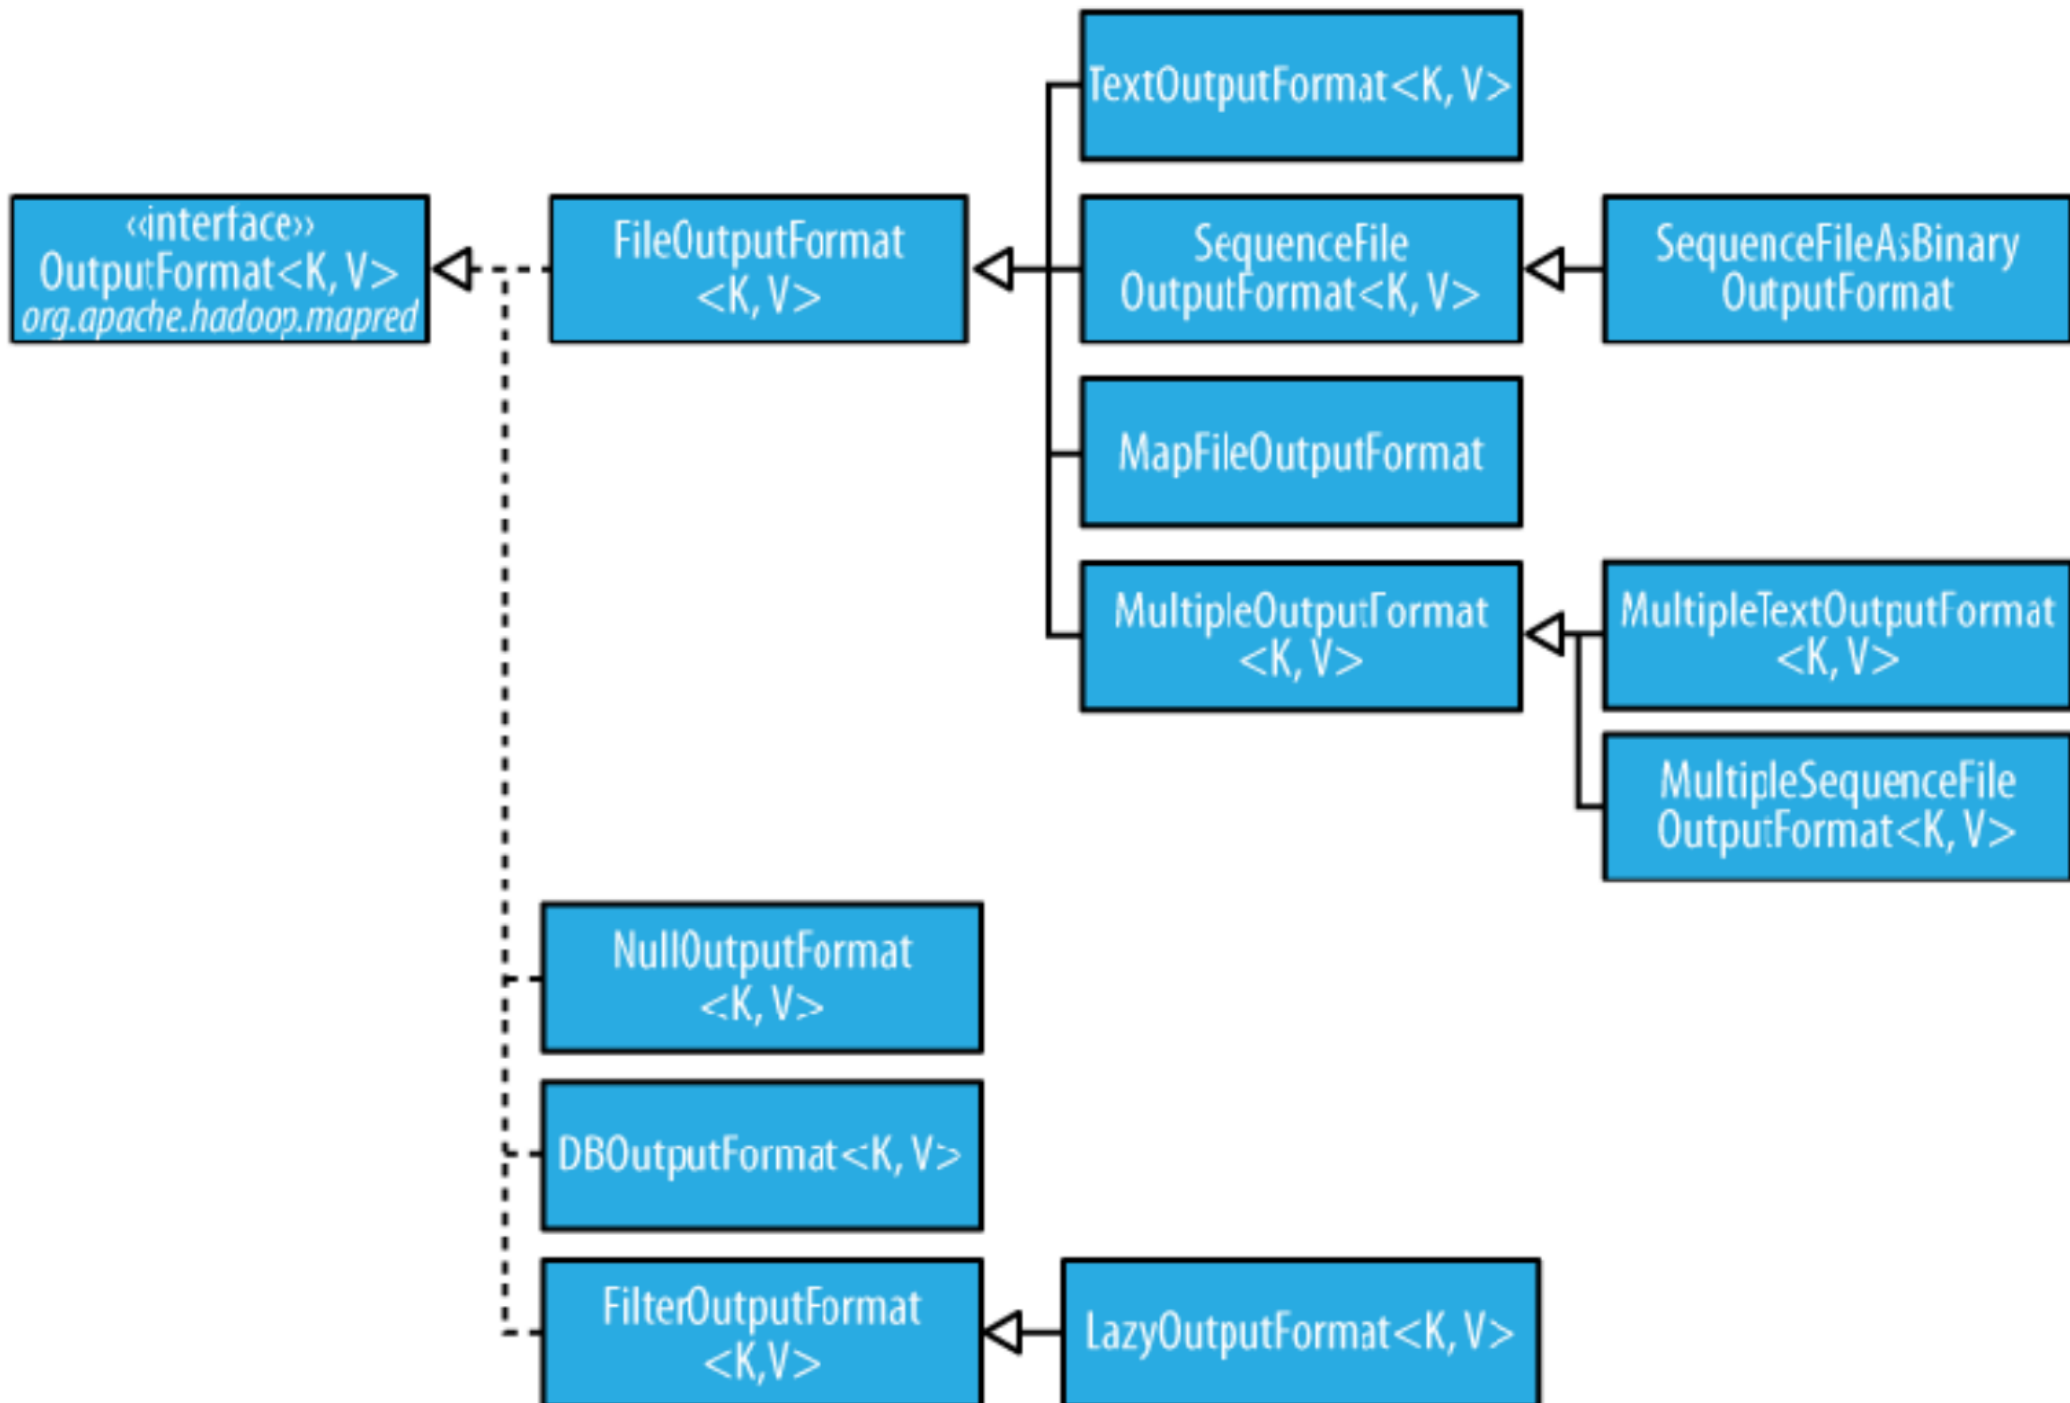


OutputFormat

Validate output specifications

RecordWriter implementation to write output files of the job

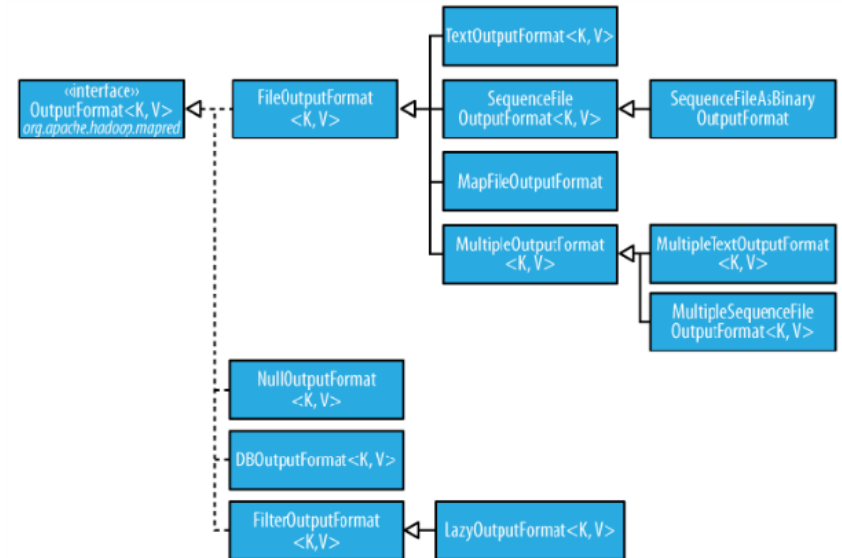




OutputFormat

Validate output specifications



RecordWriter implementation to write output files of the job

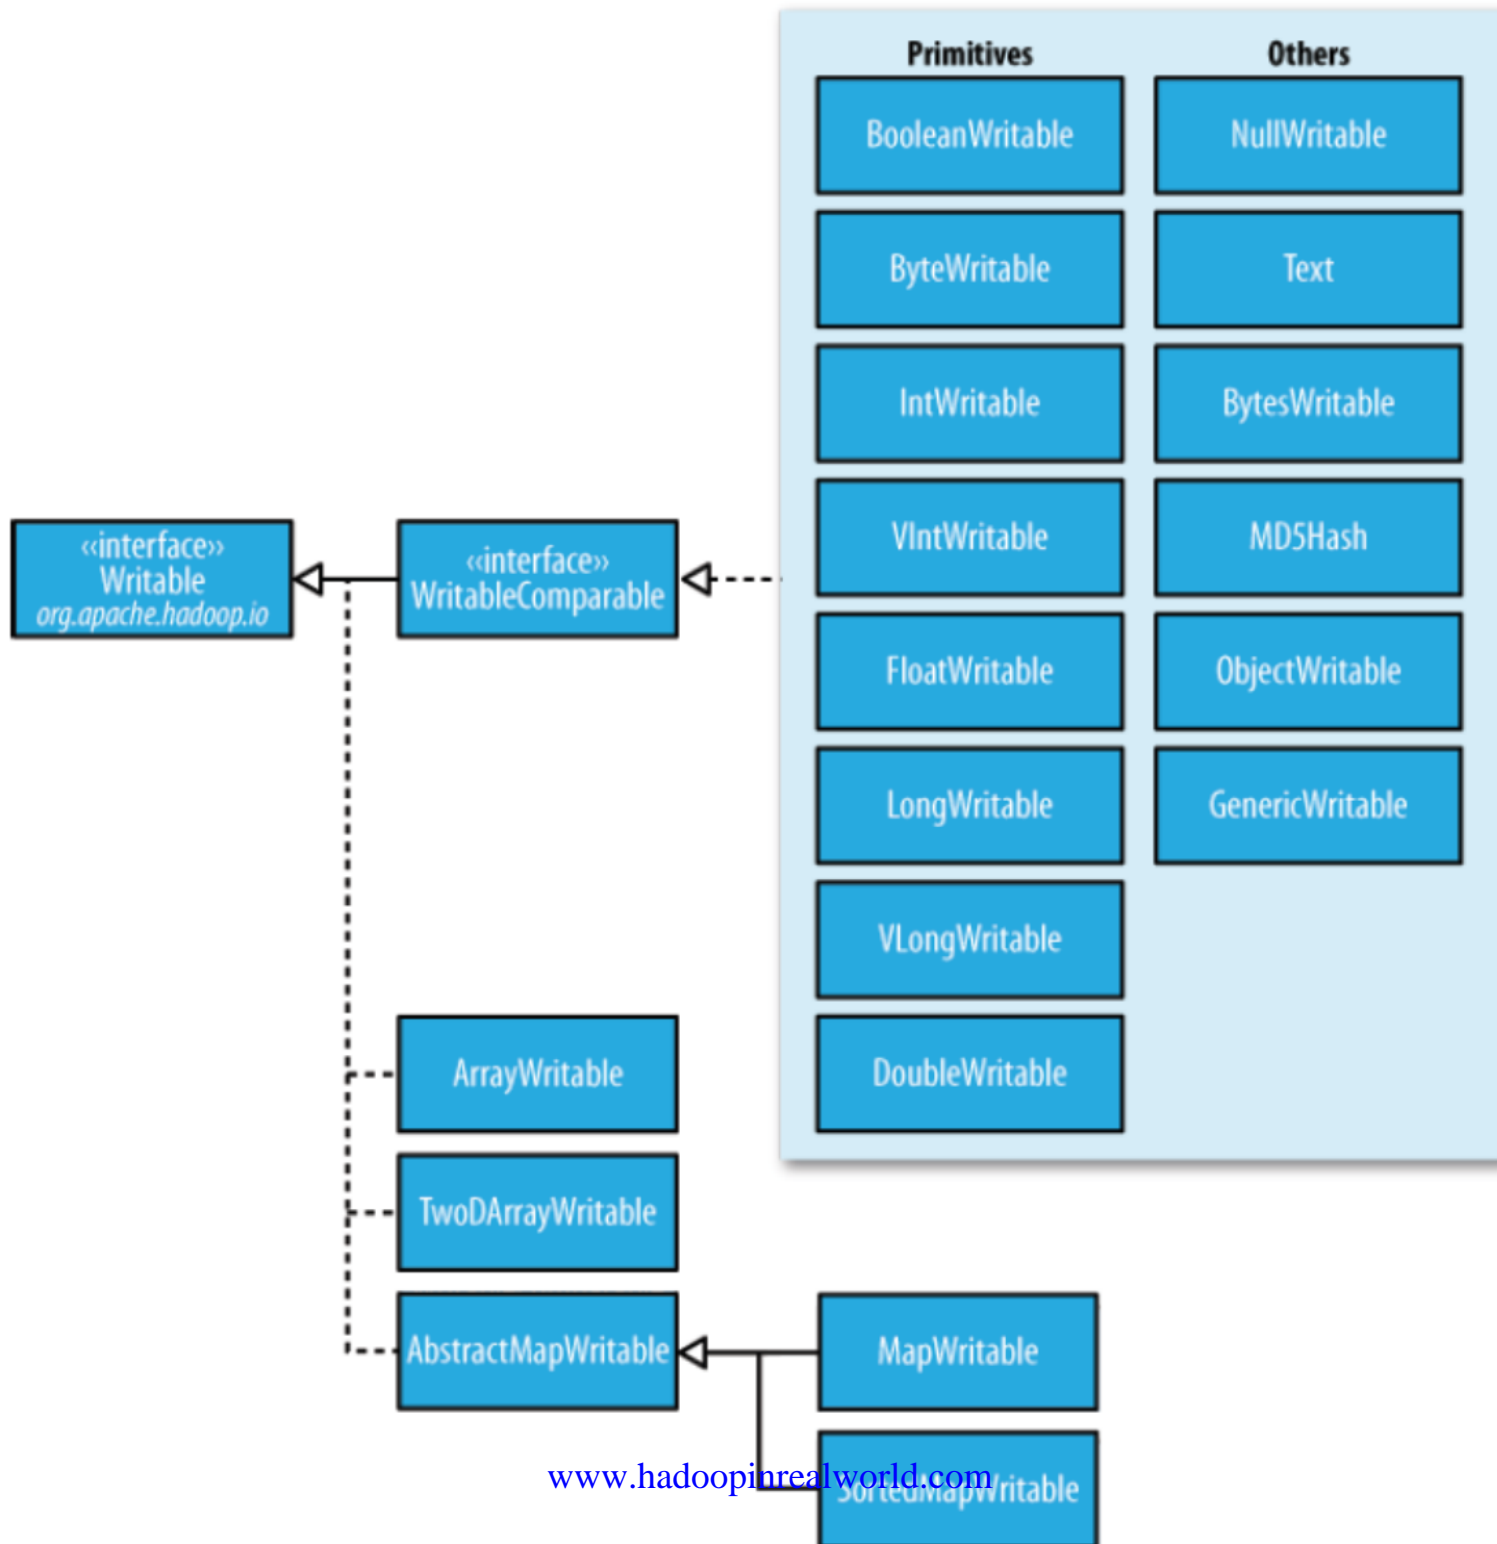



Writable

A serializable object which implements a simple, efficient, serialization protocol

Fast, compact & effective



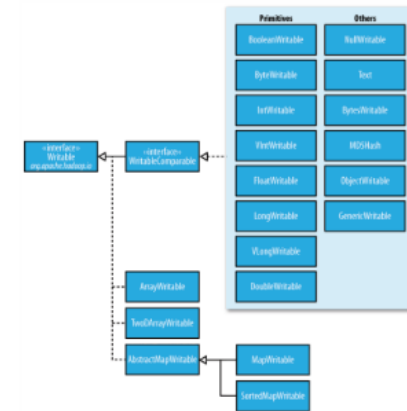





Writable

A serializable object which implements a simple, efficient, serialization protocol

Fast, compact & effective



MAPPER

- Dataset is divided in to multiple parts - Input Splits
- Each Mapper process an Input Split
- Each Mapper can be called multiple times depending on the content of Input Split
- Mapper will emit Key Value pairs as output
- There will be one or more Mapper in a MapReduce job



REDUCER

- Reduce function will take Key Value pairs from multiple Map functions as input and Reduce them to output
- Keys are grouped with values. Reduce function is called once per key and its values.
- There could be 0, 1 or more Reduce function for a MapReduce job

