

name	value	description
hadoop.hdfs.configuration.version	1	version of this configuration
dfs.namenode.rpc-address		RPC address that handles all namenodes exist, the name s address.ns1 dfs.namenode.r property will take the form c
dfs.namenode.rpc-bind-host		The actual address the RPC only the hostname portion o node or name service for HA/ all interfaces by setting it to
dfs.namenode.servicerpc-address		RPC address for HDFS Serv services should be connectir HA/Federation where multi e.g. dfs.namenode.servicerpc address.EXAMPLENAMES host1:rpc-port. If the value c will be used as the default.
dfs.namenode.servicerpc-bind-host		The actual address the servic overrides only the hostname specified per name node or 1 name node listen on all inter
dfs.namenode.lifeline.rpc-address		NameNode RPC lifeline add to isolate health checks and handler pool. In the case of 1 service ID is added to the na this property will take the fo NameNode will not start a li
dfs.namenode.lifeline.rpc-bind-host		The actual address the lifelin overrides only the hostname specified per name node or 1 name node listen on all inter
dfs.namenode.secondary.http-address	0.0.0.0:50090	The secondary namenode ht
dfs.namenode.secondary.https-address	0.0.0.0:50091	The secondary namenode H
dfs.datanode.address	0.0.0.0:50010	The datanode server address
dfs.datanode.http.address	0.0.0.0:50075	The datanode http server add
dfs.datanode.ipc.address	0.0.0.0:50020	The datanode ipc server add
dfs.datanode.http.internal-proxy.port	0	The datanode's internal web runtime.
dfs.datanode.handler.count	10	The number of server thread
dfs.namenode.http-address	0.0.0.0:50070	The address and the base po
dfs.namenode.http-bind-host		The actual address the HTTP only the hostname portion o node or name service for HA/ server listen on all interfaces
dfs.namenode.heartbeat.recheck-interval	300000	This time decides the interv dfs.heartbeat.interval, the in The unit of this configuratio
dfs.http.policy	HTTP_ONLY	Decide if HTTPS(SSL) is su HDFS daemons: The follow only on http - HTTPS_ONL Service is provided both on
dfs.client.https.need-auth	false	Whether SSL client certifica
dfs.client.cached.conn.retry	3	The number of times the HL is exceeded, the client will t
dfs.https.server.keystore.resource	ssl-server.xml	Resource file from which ss
dfs.client.https.keystore.resource	ssl-client.xml	Resource file from which ss
dfs.datanode.https.address	0.0.0.0:50475	The datanode secure http sei
dfs.namenode.https-address	0.0.0.0:50470	The namenode secure http se
dfs.namenode.https-bind-host		The actual address the HTTP only the hostname portion o node or name service for HA/ server listen on all interfaces
dfs.datanode.dns.interface	default	The name of the Network In e.g. eth2. This setting may b are assigned multiple hostna hostname. Prefer using hado
dfs.datanode.dns.nameserver	default	The host name or IP address determine its own host name dfs.datanode.dns.nameserve
dfs.namenode.backup.address	0.0.0.0:50100	The backup node server add port.

dfs.namenode.backup.http-address	0.0.0.0:50105	The backup node http server free port.
dfs.namenode.replication.considerLoad	true	Decide if choose Target cons
dfs.namenode.replication.considerLoad.factor	2.0	The factor by which a node's only if considerLoad is true.
dfs.default.chunk.view.size	32768	The number of bytes to view
dfs.datanode.du.reserved	0	Reserved space in bytes per Specific storage type based on with corresponding storage type heterogeneous storage. For a configured using property 'du reservation is not configured
dfs.namenode.name.dir	file://\${hadoop.tmp.dir}/dfs/name	Determines where on the local table(fsimage). If this is a cell replicated in all of the direct
dfs.namenode.name.dir.restore	false	Set to true to enable NameNode dfs.namenode.name.dir. When during checkpoint.
dfs.namenode.fs-limits.max-component-length	255	Defines the maximum number value of 0 will disable the cl
dfs.namenode.fs-limits.max-directory-items	1048576	Defines the maximum number property to a value less than
dfs.namenode.fs-limits.min-block-size	1048576	Minimum block size in byte accidental creation of files with performance.
dfs.namenode.fs-limits.max-blocks-per-file	1048576	Maximum number of blocks creation of extremely large f
dfs.namenode.edits.dir	\${dfs.namenode.name.dir}	Determines where on the local (edits) file. If this is a cell replicated in all of the direct dfs.namenode.name.dir
dfs.namenode.edits.dir.required		This should be a subset of directories in these places is always up-
dfs.namenode.shared.edits.dir		A directory on shared storage directory will be written by all namespaces synchronized. The dfs.namenode.edits.dir above
dfs.namenode.edits.journal-plugin.qjournal	org.apache.hadoop.hdfs.qjournal.client.QuorumJournalManager	
dfs.permissions.enabled	true	If "true", enable permission off, but all other behavior is does not change the mode, o
dfs.permissions.superusergroup	supergroup	The name of the group of su
dfs.cluster administrators		ACL for the admins, this controls servlets in the namenode, etc groups. The user list comes "user1,user2 group1,group2" "user1 group1", "user1,user2 group1"). '*' grants access to
dfs.namenode.acls.enabled	false	Set to true to enable support disabled. When ACLs are disabled getting ACLs.
dfs.namenode.lazypersist.file.scrub.interval.sec	300	The NameNode periodically and unlinks them from the namespace successive scans. Set it to a
dfs.block.access.token.enable	false	If "true", access tokens are used tokens are checked on acces
dfs.block.access.key.update.interval	600	Interval in minutes at which
dfs.block.access.token.lifetime	600	The lifetime of access token
dfs.datanode.data.dir	file://\${hadoop.tmp.dir}/dfs/data	Determines where on the local is a comma-delimited list of typically on different device types ([SSD]/[DISK]/[ARC] storage type will be DISK if Directories that do not exist
dfs.datanode.data.dir.perm	700	Permissions for the directory blocks. The permissions can
dfs.replication	3	Default block replication. The is created. The default is use
dfs.replication.max	512	Maximal block replication.
dfs.namenode.replication.min	1	Minimal block replication.
dfs.namenode.maintenance.replication.min	1	Minimal live block replicati
dfs.namenode.safemode.replication.min		a separate minimum replicat

3/17

dfs.namenode.service.handler.count	10	The number of Namenode R from all other non-client noc dfs.namenode.servicerpc-ad
dfs.namenode.lifeline.handler.ratio	0.10	A ratio applied to the value of number of RPC server thread. For example, if dfs.namenode.lifeline.handler.ratio is 0.10, then the NameNode starts 10% of the number of RPC server threads. It is common to tune the number of DataNodes in a cluster to be the same as the number of lifeline handler threads to be tuned. If the value of dfs.namenode.lifeline.handler.ratio is less than 0, then the property has no effect if dfs.namenode.lifeline.handler.count is defined.
dfs.namenode.lifeline.handler.count		Sets an absolute number of lifeline DataNode Lifeline Protocol handlers. If defined, then it overrides the value of dfs.namenode.lifeline.handler.ratio. If not defined, this property has no effect.
dfs.namenode.safemode.threshold-pct	0.999f	Specifies the percentage of total blocks defined by dfs.namenode.replication.min for any particular percentage of blocks to make safe mode permanent.
dfs.namenode.safemode.min.datanodes	0	Specifies the number of data nodes in safe mode. Values less than 0 are ignored. account when deciding whether to leave the number of datanodes in the cluster.
dfs.namenode.safemode.extension	30000	Determines extension of safe mode. Support multiple time unit suffixes.
dfs.namenode.resource.check.interval	5000	The interval in milliseconds between calculations of the number of the blocks. If the number of blocks is less than dfs.namenode.resource.check.minimum, then the minimum value specified.
dfs.namenode.resource.du.reserved	104857600	The amount of space to reserve for local edits. The default is 100MB.
dfs.namenode.resource.checked.volumes		A list of local directories for local edits directories.
dfs.namenode.resource.checked.volumes.minimum	1	The minimum number of required volumes.
dfs.datanode.balance.bandwidthPerSec	10m	Specifies the maximum amount of bandwidth for balancing purpose in terms of bytes per second. Suffix (case insensitive): k(kilobyte), m(megabyte), g(gigabyte) (such as 128k, 512m, 1g, etc). The default is 128 MB).
dfs.hosts		Names a file that contains a list of hostnames. The full pathname of the file must be specified.
dfs.hosts.exclude		Names a file that contains a list of hostnames to be excluded. The full pathname of the file must be specified.
dfs.namenode.max.objects	0	The maximum number of files in the namespace. If 0, indicates no limit to the number of files.
dfs.namenode.datanode.registration.ip-hostname-check	true	If true (the default), then the IP address must be resolved to a hostname. If false, the IP address can be left on to prevent the NameNode from excluding file during a DNS infrastructure to support reverse lookup.
dfs.namenode.decommission.interval	30	Namenode periodicity in seconds to check for decommission. Support multiple time unit suffixes.
dfs.namenode.decommission.blocks.per.interval	500000	The approximate number of blocks to be decommissioned as defined in dfs.namenode.decommission.interval.
dfs.namenode.decommission.max.concurrent.tracked.nodes	100	The maximum number of decommissioned nodes that will be tracked at any time. If the number of nodes is greater than this conservative limit, the number of nodes at once is reduced.
dfs.namenode.replication.interval	3	The periodicity in seconds to check for replication. Support multiple time unit suffixes.
dfs.namenode.access.time.precision	3600000	The access time for HDFS files. Setting a value of 0 disables access time.
dfs.datanode.plugins		Comma-separated list of datanode plugins.
dfs.namenode.plugins		Comma-separated list of namenode plugins.
dfs.namenode.block-placement-policy.default.prefer-local-node	true	Controls how the default block placement policy behaves. If true, it will prefer the node with the same rack as the client. If false, it will prefer the node with the same rack as the server. If the file ends up on a single node, the policy has no effect.

dfs.stream-buffer-size	4096	The size of buffer to stream hardware page size (4096 or read and write operations.
dfs.bytes-per-checksum	512	The number of bytes per che
dfs.client-write-packet-size	65536	Packet size for clients to wri
dfs.client.write.exclude.nodes.cache.expiry.interval.millis	600000	The maximum period to kee in milliseconds, the previous cache and will be considere situations where you keep a (WAL) file) to make the wri minutes.
dfs.namenode.checkpoint.dir	file://\${hadoop.tmp.dir}/dfs/namesecondary	Determines where on the loc temporary images to merge. replicated in all of the direct
dfs.namenode.checkpoint.edits.dir	\${dfs.namenode.checkpoint.dir}	Determines where on the loc temporary edits to merge. If replicated in all of the direct dfs.namenode.checkpoint.di
dfs.namenode.checkpoint.period	3600	The number of seconds betw
dfs.namenode.checkpoint.txns	1000000	The Secondary NameNode c every 'dfs.namenode.checkp 'dfs.namenode.checkpoint,pe
dfs.namenode.checkpoint.check.period	60	The SecondaryNameNode a 'dfs.namenode.checkpoint.cl transactions.
dfs.namenode.checkpoint.max-retries	3	The SecondaryNameNode r fsimage or replaying edits, t
dfs.namenode.num.checkpoints.retained	2	The number of image check and Secondary NameNode i necessary to recover an up-t be retained.
dfs.namenode.num.extra.edits.retained	1000000	The number of extra transac necessary for a NN restart. I kept, but to the number of tr contain several transactions number of edits to retain as subtracted by the value of th that don't include the compu and purges them subsequent HA setup where a remote St have a longer backlog of ret order of a few hundred byte hundreds of MBs or low GB specified for this setting if d the number configured by di
dfs.namenode.max.extra.edits.segments.retained	10000	The maximum number of ex minimally necessary for a N dfs.namenode.num.extra.edi number of extra edits files to
dfs.namenode.delegation.key.update-interval	86400000	The update interval for mast
dfs.namenode.delegation.token.max-lifetime	604800000	The maximum lifetime in m
dfs.namenode.delegation.token.renew-interval	86400000	The renewal interval for del
dfs.datanode.failed.volumes.tolerated	0	The number of volumes that default any volume failure w
dfs.image.compress	false	Should the dfs image be con
dfs.image.compression.codec	org.apache.hadoop.io.compress.DefaultCodec	If the dfs image is compress defined in io.compression.cc
dfs.image.transfer.timeout	60000	Socket timeout for image tra dfs.image.transfer.bandwidth image transfer can complete sender fails during image tra
dfs.image.transfer.bandwidthPerSec	0	Maximum bandwidth used f namenode), in bytes per sec responsive during checkpoin dfs.image.transfer.timeout sl successfully. A default value bandwidth used for bootstrap bootstrap-standby.bandwidthl
dfs.image.transfer-bootstrap-standby.bandwidthPerSec	0	Maximum bandwidth used f per second. A default value should be used in most case used for regular image trans
dfs.image.transfer.chunksize	65536	Chunksize in bytes to uploa buffering of contents of ima

dfs.namenode.support.allow.format	true	Does HDFS namenode allow for any production cluster, t
dfs.datanode.max.transfer.threads	4096	Specifies the maximum num
dfs.datanode.scan.period.hours	504	If this is positive, the DataN specified scan period. If this then the default value of 504 documented that setting this
dfs.block.scanner.volume.bytes.per.second	1048576	If this is 0, the DataNode's b number of bytes per second volume.
dfs.datanode.readahead.bytes	4194304	While reading block files, if the posix_fadvise system cal ahead of the current reader's are highly contended. This c read position which the data by configuring this property has no effect.
dfs.datanode.drop.cache.behind.reads	false	In some workloads, the data it is unlikely to be useful to , DataNode may be configure is delivered to the client. Th only short sections of a bloc performance for some workl data. If the Hadoop native li
dfs.datanode.drop.cache.behind.writes	false	In some workloads, the data it is unlikely to be useful to , DataNode may be configure is written to disk. This may : cache space usage for more this configuration has no eff
dfs.datanode.sync.behind.writes	false	If this configuration is enabl all written data to the disk in policy which may wait for u performance for some workl Hadoop native libraries are 1
dfs.client.failover.max.attempts	15	Expert only. The number of is considered failed.
dfs.client.failover.sleep.base.millis	500	Expert only. The time to wai exponentially as a function of +/- 50%. This option specifi failover will retry immediate dfs.client.failover.sleep.base
dfs.client.failover.sleep.max.millis	15000	Expert only. The time to wai exponentially as a function of +/- 50%. This option specifi the time between two failove dfs.client.failover.sleep.max
dfs.client.failover.connection.retries	0	Expert only. Indicates the nu server connection.
dfs.client.failover.connection.retries.on.timeouts	0	Expert only. The number of when establishing a server c
dfs.client.datanode-restart.timeout	30	Expert only. The time to wai notification for quick restart recovery mechanisms. The r using the shutdownDatanode
dfs.nameservices		Comma-separated list of nar
dfs.nameservice.id		The ID of this nameservice. nameservice is configured fo the local node's address with
dfs.internal.nameservices		Comma-separated list of nar all the nameservices in this l
dfs.ha.namenodes.EXAMPLENAMESERVICE		The prefix for a given name: given nameservice (eg EXA NameNode in the nameservi determine all the NameNode nameservice ID previously, : the NameNodes, you would "nn1,nn2".
dfs.ha.namenode.id		The ID of this namenode. If automatically by matching t
dfs.ha.log-roll.period	120	How often, in seconds, the S StandbyNode only reads fro to-date as how often the log: StandbyNode will be up to c
dfs.ha.tail-edits.period	60	How often, in seconds, the S

		the shared edits log.
dfs.ha.tail-edits.rolleddits.timeout	60	The timeout in seconds of ca
dfs.ha.automatic-failover.enabled	false	Whether automatic failover details on automatic HA cor
dfs.client.use.datanode.hostname	false	Whether clients should use c
dfs.datanode.use.datanode.hostname	false	Whether datanodes should u data transfer.
dfs.client.local.interfaces		A comma separated list of n client and datanodes. When client chooses one of the spe interface. Individual names i subinterface name (eg "eth0 notation to match a range of
dfs.datanode.shared.file.descriptor.paths	/dev/shm,/tmp	A comma-separated list of p between the DataNode and t descriptors will not be writte /tmp by default.
dfs.short.circuit.shared.memory.watcher.interrupt.check.ms	60000	The length of time in millise between checking for java in for unit tests.
dfs.namenode.kerberos.principal		The NameNode service prin NameNode will substitute _HOST placeholder allows i HA setup.
dfs.namenode.keytab.file		The keytab file used by each principal name is configurec
dfs.datanode.kerberos.principal		The DataNode service princ DataNode will substitute _H HOST placeholder allows i
dfs.datanode.keytab.file		The keytab file used by each principal name is configurec
dfs.journalnode.kerberos.principal		The JournalNode service pri JournalNode will substitute _HOST placeholder allows i
dfs.journalnode.keytab.file		The keytab file used by each principal name is configurec
dfs.namenode.kerberos.internal.spnego.principal	\${dfs.web.authentication.kerberos.principal}	The server principal used by Kerberos security is enabled SPNEGO server principal be the web server will attempt t dfs.web.authentication.kerbe
dfs.journalnode.kerberos.internal.spnego.principal		The server principal used by when Kerberos security is en The SPNEGO server princip "*", the web server will atten dfs.web.authentication.kerbe \${dfs.web.authentication.ke dfs.web.authentication.kerbe
dfs.secondary.namenode.kerberos.internal.spnego.principal	\${dfs.web.authentication.kerberos.principal}	The server principal used by authentication when Kerbero settings, it is ignored in an E with every principal specifie
dfs.web.authentication.kerberos.principal		The server principal used by Required when WebHDFS a also used to specify the valu dfs.journalnode.kerberos.int
dfs.web.authentication.kerberos.keytab		The keytab file for the princ
dfs.namenode.kerberos.principal.pattern	*	A client-side RegEx that car (useful in cross-realm env.)
dfs.namenode.avoid.read.stale.datanode	false	Indicate whether or not to av have not been received by th datanodes will be moved to dfs.namenode.avoid.write.sta
dfs.namenode.avoid.write.stale.datanode	false	Indicate whether or not to av have not been received by th avoid using stale datanodes i (dfs.namenode.write.stale.da dfs.namenode.avoid.read.sta
dfs.namenode.stale.datanode.interval	30000	Default time interval in mill namenode has not received l the datanode will be marked too small since otherwise thi minimum stale interval valu guarantee that the stale inter

		avoided during lease/block r dfs.namenode.avoid.read.sta dfs.namenode.avoid.write.st
dfs.namenode.write.stale.datanode.ratio	0.5f	When the ratio of number st ratio, stop avoiding writing t
dfs.namenode.invalidate.work.pct.per.iteration	0.32f	*Note*: Advanced property. of block invalidations (delet final deletion count is deter in the system. The resultant for proper invalidation over non-zero percentage in float
dfs.namenode.replication.work.multiplier.per.iteration	2	*Note*: Advanced property. block transfers to begin in p being sent over a DN heartb this multiplier with the total number of blocks to begin tr any positive, non-zero integ
nfs.server.port	2049	Specify the port number use
nfs.mountd.port	4242	Specify the port number use
nfs.dump.dir	/tmp/.hdfs-nfs	This directory is used to tem each file, the out-of-order w threshold (e.g., 1MB) in me
nfs.rtxmax	1048576	This is the maximum size in you change this, make sure the mount directive).
nfs.wtmax	1048576	This is the maximum size in you change this, make sure the mount directive).
nfs.keytab.file		*Note*: Advanced property. hdfs-nfs gateway. This is rec
nfs.kerberos.principal		*Note*: Advanced property. principal. This is required w gateway-user/nfs-gateway-h
nfs.allow.insecure.ports	true	When set to false, client con 1023) will be rejected. This have had root privilege on th
dfs.webhdfs.enabled	true	Enable WebHDFS (REST A
hadoop.fuse.connection.timeout	300	The minimum number of sec Lower values will result in l by avoiding the overhead of
hadoop.fuse.timer.period	5	The number of seconds betw in fuse dfs noticing changes
dfs.namenode.metrics.logger.period.seconds	600	This setting controls how fre configuration must also defi metrics to be logged. Name less than zero.
dfs.datanode.metrics.logger.period.seconds	600	This setting controls how fre configuration must also defi metrics to be logged. DataN than zero.
dfs.metrics.percentiles.intervals		Comma-delimited set of inte percentile latency metrics or metrics are disabled.
dfs.datanode.peer.stats.enabled	false	A switch to turn on/off track
dfs.datanode.outliers.report.interval	1800000	This setting controls how fre NameNode via heartbeats. T dfs.heartbeat.interval. If no dfs.datanode.peer.stats.enabl
dfs.datanode.fileio.profiling.sampling.percentage	0	This setting controls the per disk statistics. The default v 1 and 100 to enable disk stat
hadoop.user.group.metrics.percentiles.intervals		A comma-separated list of tl 50/75/90/95/99th percentile percentile latency metrics ar
dfs.encrypt.data.transfer	false	Whether or not actual block the wire. This only needs to automatically. It is possible logic via dfs.trustedchannel.
dfs.encrypt.data.transfer.algorithm		This value may be set to eith default on the system is usec cryptographically secure, bu both the client and server th transfer keys for AES. (See



dfs.encrypt.data.transfer.cipher.suites		This value may be either <code>dfs.encrypt.data.transfer</code> use then only the algorithm specified the property is not defined.
dfs.encrypt.data.transfer.cipher.key.bitlength	128	The key bitlength negotiated set to either 128, 192 or 256
dfs.trustedchannel.resolver.class		TrustedChannelResolver is a transfer. The TrustedChannel resolver indicates that the channel is even if <code>dfs.encrypt.data.transfer</code> indicating that the channel is
dfs.data.transfer.protection		A comma-separated list of SecurityDataNode when reading or writing and privacy. authentication and integrity and privacy are enabled supersedes the setting for <code>dfs.data.transfer.protection</code> use a specialized encrypted DataNode listening on a privileged port establishes security
dfs.data.transfer.saslproperties.resolver.class		SaslPropertiesResolver used when reading or writing blocks <code>hadoop.security.saslproperties</code>
dfs.datanode.hdfs-blocks-metadata.enabled	false	Boolean which enables back DistributedFileSystem#getF
dfs.client.file-block-storage-locations.num-threads	10	Number of threads used for DistributedFileSystem#getF
dfs.client.file-block-storage-locations.timeout.millis	1000	Timeout (in milliseconds) for DistributedFileSystem#getF
dfs.journalnode.rpc-address	0.0.0.0:8485	The JournalNode RPC server
dfs.journalnode.http-address	0.0.0.0:8480	The address and port the Journal server will start on a free port
dfs.journalnode.https-address	0.0.0.0:8481	The address and port the Journal server will start on a free port
dfs.namenode.audit.loggers	default	List of classes implementing implementations of <code>org.apache.hadoop.log</code> value "default" can be used to log system. Installing custom NameNode. Refer to the custom
dfs.datanode.available-space-volume-choosing-policy.balanced-space-threshold	10737418240	Only used when the <code>dfs.data.org.apache.hadoop.hdfs.service</code> This setting controls how much disk space before they are considered within this range of each other assignments will be done on
dfs.datanode.available-space-volume-choosing-policy.balanced-space-preference-fraction	0.75f	Only used when the <code>dfs.data.org.apache.hadoop.hdfs.service</code> This setting controls what percentage more available disk space than in practice 0.5 - 1.0, since the available disk space receive
dfs.namenode.edits.noeditlogchannelflush	false	Specifies whether to flush edit calls are skipped and synchronize file with RandomAccessFile of edit log writes on the Write platform and hardware specific FileChannel#force. For exabyte devices while it might not be a caution.
dfs.client.cache.drop.behind.writes		Just like <code>dfs.datanode.drop.cache</code> dropped behind HDFS write <code>dfs.datanode.drop.cache.behind</code> the entire datanode. If present libraries are not available to
dfs.client.cache.drop.behind.reads		Just like <code>dfs.datanode.drop.cache</code> dropped behind HDFS reads <code>dfs.datanode.drop.cache.behind</code> the entire datanode. If present libraries are not available to
dfs.client.cache.readahead		When using remote reads, the using <code>posix_fadvise</code> , potentially <code>dfs.datanode.readahead.byte</code> datanode. If present, this setting determines reads, this setting determines native libraries are not available

dfs.client.server-defaults.validity.period.ms	3600000	The amount of milliseconds parameter is set to 1 hour.
dfs.namenode.enable.retrycache	true	This enables the retry cache requests the corresponding r retry cache is sent. Such ope protocols. It is recommende clients getting failure respon for transparent fail-overs. Tl dfs.namenode.retrycache.ex
dfs.namenode.retrycache.expirytime.millis	600000	The time for which retry cac
dfs.namenode.retrycache.heap.percent	0.03f	This parameter configures th cached). This corresponds to process java heap size. Assu dfs.namenode.retrycache.ex support 7 operations per sec operation rate linearly increa
dfs.client.mmap.enabled	true	If this is set to false, the clie
dfs.client.mmap.cache.size	256	When zero-copy reads are u mapped regions. This param in that cache. The larger this memory-mapped files. mma increase your ulimit virtual s size. Note that you can still c
dfs.client.mmap.cache.timeout.ms	3600000	The minimum length of time an entry is in the cache long background thread.
dfs.client.mmap.retry.timeout.ms	300000	The minimum amount of tin
dfs.client.short.circuit.replica.stale.threshold.ms	1800000	The maximum amount of tir there is no communication f the short-circuit replica ever
dfs.namenode.path.based.cache.block.map.allocation.percent	0.25	The percentage of the Java b cached blocks map is a hash accessed more slowly if the more memory.
dfs.datanode.max.locked.memory	0	The amount of memory in b datanode. The datanode's m must be set to at least this va parameter is set to 0, which available to the DataNode, ti
dfs.namenode.list.cache.directives.num.responses	100	This value controls the numl wire in response to a listDir
dfs.namenode.list.cache.pools.num.responses	100	This value controls the numl in response to a listPools RP
dfs.namenode.path.based.cache.refresh.interval.ms	30000	The amount of milliseconds when we calculate which blo parameter is set to 30 secon
dfs.namenode.path.based.cache.retry.interval.ms	30000	When the NameNode needs not cached, it must direct the DNA_UNCACHE comman how frequently the NameNo
dfs.datanode.fsdatasetcache.max.threads.per.volume	4	The maximum number of th These threads consume both
dfs.cachereport.intervalMsec	10000	Determines cache reporting DataNode sends a full repor cache report to update its m has no effect if in-memory c dfs.datanode.max.locked.me available to the DataNode, ti
dfs.namenode.edit.log.autoroll.multiplier.threshold	2.0	Determines when an active i number of edits) is determin This prevents extremely larg can cause timeouts during n behavior is intended as a fail the edit log by the normal ch
dfs.namenode.edit.log.autoroll.check.interval.ms	300000	How often an active namenc
dfs.webhdfs.user.provider.user.pattern	^[A-Za-z_][A-Za-z0-9._-]*[\$]?\$	Valid pattern for user and gr
dfs.webhdfs.acl.provider.permission.pattern	^(default:)?(user group mask other):[[A-Za-z_][A-Za-z0-9._-]]*:([rwx-]{3})?((default:)?(user group mask other):[[A-Za-z_][A-Za-z0-9._-]]*:([rwx-]{3})?)*\$	Valid pattern for user and gr regex.
dfs.webhdfs.socket.connect-timeout	60s	Socket timeout for connecti hanging if the server hostna timeout expires. Value is fol microseconds, milliseconds, provide units, but millisecor
dfs.webhdfs.socket.read-timeout	60s	Socket timeout for reading d

		from hanging if the server st ms, s, m, h, d for nanosecon respectively. Values should j
dfs.client.context	default	The name of the DFSCli socket cache and short-circu you don't want to share with
dfs.client.read.shortcircuit	false	This configuration paramete
dfs.client.socket.send.buffer.size	0	Socket send buffer size for a connection throughput. If it explicitly, thus enable tcp au
dfs.domain.socket.path		Optional. This is a path to a between the DataNode and l path, it will be replaced by tl
dfs.client.read.shortcircuit.skip.checksum	false	If this configuration paramet normally not recommended, using this if you are doing y
dfs.client.read.shortcircuit.streams.cache.size	256	The DFSCli controls the maximum numb more file descriptors, but po lots of seeks.
dfs.client.read.shortcircuit.streams.cache.expiry.ms	300000	This controls the minimum a context before they can be c
dfs.datanode.shared.file.descriptor.paths	/dev/shm,/tmp	Comma separated paths to tl The client and the DataNode paths in order until creation
dfs.namenode.audit.log.debug.cmdlist		A comma separated list of N audit log only if the audit lo
dfs.client.use.legacy.blockreader.local	false	Legacy short-circuit reader i configuration parameter is tr implementation based on HI
dfs.block.local-path-access.user		Comma separated list of the read.
dfs.client.domain.socket.data.traffic	false	This control whether we wil rather than over TCP socket turned off by default.
dfs.namenode.reject-unresolved-dn-topology-mapping	false	If the value is set to true, the mapping for a datanode is no net.topology.script.file.name the default rack will be assign data resiliency, since they de datanode registration with th
dfs.client.slow.io.warning.threshold.ms	30000	The threshold in millisecond default, this parameter is set
dfs.datanode.slow.io.warning.threshold.ms	300	The threshold in millisecond default, this parameter is set
dfs.namenode.xattr.enabled	true	Whether support for extende
dfs.namenode.fs-limits.max-xattrs-per-inode	32	Maximum number of extenc
dfs.namenode.fs-limits.max-xattr-size	16384	The maximum combined siz should be larger than 0, and
dfs.namenode.lease-recheck-interval-ms	2000	During the release of lease a In order to not block them d lock limit.
dfs.namenode.max-lock-hold-to-release-lease-ms	25	During the release of lease a In order to not block them d lock limit.
dfs.namenode.write-lock-reporting-threshold-ms	5000	When a write lock is held or is released. This sets how lo
dfs.namenode.read-lock-reporting-threshold-ms	5000	When a read lock is held on released. This sets how long
dfs.namenode.lock.detailed-metrics.enabled	false	If true, the namenode will ke lock for and emit this as met FSN(Read Write)LockNano the operation that initiated th operations) and they export t
dfs.namenode.fslock.fair	true	If this is true, the FS Names prevent writer threads from i java.util.concurrent.locks.Re locks.
dfs.namenode.startup.delay.block.deletion.sec	0	The delay in seconds at whic By default it's disabled. In th deleted, suggested delay is c number of pending deletion
dfs.namenode.list.encryption.zones.num.responses	100	When listing encryption zon

		batch. Fetching the list incre
dfs.namenode.list.openfiles.num.responses	1000	When listing open files, the batch. Fetching the list incre
dfs.namenode.edekcacheloader.interval.ms	1000	When KeyProvider is config up / becomes active. All ede cache loader will try to warr
dfs.namenode.edekcacheloader.initial.delay.ms	3000	When KeyProvider is config cache on NN start up / becoi
dfs.namenode.inotify.max.events.per.rpc	1000	Maximum number of events The default value attempts to huge memory requirements more than 1 MB.)
dfs.user.home.dir.prefix	/user	The directory to prepend to
dfs.datanode.cache.revocation.timeout.ms	900000	When the DFSCClient reads f can skip verifying checksum client is done. If the client ta evict the block file from the wait for the client to release
dfs.datanode.cache.revocation.polling.ms	500	How often the DataNode sho the DataNode wants to unca
dfs.datanode.block.id.layout.upgrade.threads	12	The number of threads to us during upgrade of a DataNo on the layout).
dfs.storage.policy.enabled	true	Allow users to change the st
dfs.namenode.legacy-oiv-image.dir		Determines where to save th by standby NameNode or Sc format fsimage by oiv_legacy will not be saved in checkpc
dfs.namenode.top.enabled	true	Enable nntop: reporting top
dfs.namenode.top.window.num.buckets	10	Number of buckets in the ro
dfs.namenode.top.num.users	10	Number of top users returne
dfs.namenode.top.windows.minutes	1,5,25	comma separated list of nntc
dfs.webhdfs.ugi.expire.after.access	600000	How long in milliseconds af expire.
dfs.namenode.blocks.per.postponedblocks.rescan	10000	Number of blocks to rescan
dfs.datanode.block-pinning.enabled	false	Whether pin blocks on favor
dfs.client.block.write.locateFollowingBlock.initial.delay.ms	400	The initial delay (unit is ms) exponentially(double) for ea
dfs.ha.zkfc.nn.http.timeout.ms	20000	The HTTP connection and r local NN thread dump after SERVICE_UNHEALTHY. I
dfs.namenode.quota.init-threads	4	The number of concurrent th initialization also affects the try increasing this.
dfs.datanode.transfer.socket.send.buffer.size	0	Socket send buffer size for I This may affect TCP connec size will be set explicitly, th 0.
dfs.datanode.transfer.socket.recv.buffer.size	0	Socket receive buffer size fo writing). This may affect TC no buffer size will be set ex default value is 0.
dfs.namenode.upgrade.domain.factor	\${dfs.replication}	This is valid only when bloc BlockPlacementPolicyWithl domains any block's replicas this value, the policy ensure of replicas is greater than thi least this value.
dfs.ha.zkfc.port	8019	RPC port for Zookeeper Fai
dfs.datanode.bp-ready.timeout	20	The maximum wait time for Setting this to 0 fails request namenode. This wait time re
dfs.datanode.cached-dfsused.check.interval.ms	600000	The interval check time of l doing the rolling upgrade op expired and redo the du oper Adjust this property can mal
dfs.webhdfs.rest-csrf.enabled	false	If true, then enables WebHD WebHDFS client also uses tl custom CSRF prevention he
dfs.webhdfs.rest-csrf.custom-header	X-XSRF-HEADER	The name of a custom head site request forgery (CSRF)

		csrf.enabled to true. The WebHDFS client does not need to send the csrf token.
dfs.webhdfs.rest.csrf.methods-to-ignore	GET,OPTIONS,HEAD,TRACE	A comma-separated list of HTTP methods that WebHDFS will ignore when determining if a request is a CSRF attempt. By default, WebHDFS will ignore GET, OPTIONS, HEAD, and TRACE. To protect WebHDFS by setting dfs.webhdfs.rest.csrf.methods-to-ignore to a value that matches any of these regular expressions, WebHDFS will not match any of these regular expressions, and therefore CSRF protection will not be enabled. To match any of these regular expressions, WebHDFS will match any of these regular expressions, and therefore CSRF protection will not be enabled. To match any of these regular expressions, WebHDFS will match any of these regular expressions, and therefore CSRF protection will not be enabled.
dfs.webhdfs.rest.csrf.browser-useragents-regex	^Mozilla.*;^Opera.*	A comma-separated list of regular expressions that WebHDFS will use to determine if a request is a CSRF attempt. By default, WebHDFS will use a regular expression that matches any of these regular expressions, and therefore CSRF protection will not be enabled. To protect WebHDFS by setting dfs.webhdfs.rest.csrf.browser-useragents-regex to a value that matches any of these regular expressions, WebHDFS will not match any of these regular expressions, and therefore CSRF protection will not be enabled. To match any of these regular expressions, WebHDFS will match any of these regular expressions, and therefore CSRF protection will not be enabled.
dfs.xframe.enabled	true	If true, then enables protection against cross-site scripting (XSS) attacks. The header value set to SAMEORIGIN. If false, then the header value is set to transparent or opaque layers page.
dfs.xframe.value	SAMEORIGIN	This configuration value allows you to specify the value of the X-Frame-Options header. Possible values for this field are SAMEORIGIN, DENY, and ALLOW-FROM. If the value is not one of these, an exception will be thrown.
dfs.http.client.retry.policy.enabled	false	If "true", enable the retry policy. Enabling the retry policy can cause a client to retry a failed request between clusters that could be successful during the copy.
dfs.http.client.retry.policy.spec	10000,6,60000,10	Specify a policy of multiple retries and sleep time between retries. The format is: retries,sleep time on average, the sleep time on failure.
dfs.http.client.failover.max.attempts	15	Specify the max number of retries for a failed request.
dfs.http.client.retry.max.attempts	10	Specify the max number of retries for a failed request. If the number of retries exceeds this value, there will be no more retries.
dfs.http.client.failover.sleep.base.millis	500	Specify the base amount of time to sleep between retries on failure.
dfs.http.client.failover.sleep.max.millis	15000	Specify the upper bound of time to sleep between retries on failure.
dfs.namenode.hosts.provider.classname	org.apache.hadoop.hdfs.server.blockmanagement.HostFileManager	The class that provides access to the host information for the namenode. The class org.apache.hadoop.hdfs.server.blockmanagement.HostFileManager which loads files specified by the namenode configuration file org.apache.hadoop.hdfs.server.namenode.NAMENODE_CONFIG will load the JSON file defined in the namenode configuration file "dfsadmin -refreshNodes" on the namenode.
datanode.https.port	50475	HTTPS port for DataNode.
dfs.balancer.dispatcherThreads	200	Size of the thread pool for the balancer dispatcher.
dfs.balancer.movedWinWidth	5400000	Window of time in ms for the balancer to move blocks.
dfs.balancer.moverThreads	1000	Thread pool size for executing balancer moves.
dfs.balancer.max-size-to-move	10737418240	Maximum number of bytes to move in a single operation.
dfs.balancer.getBlocks.min-block-size	10485760	Minimum block threshold size for the balancer to move.
dfs.balancer.getBlocks.size	2147483648	Total size in bytes of DataNodes to move.
dfs.balancer.block-move.timeout	0	Maximum amount of time in seconds that the balancer will stop waiting for a block to be moved. A 3 to 5 minute timeout is reasonable. If the number of moves is large, this needs to be increased. If many nodes are constantly being moved, balancer parameters might need to be adjusted.
dfs.balancer.max-no-move-interval	60000	If this specified amount of time has passed since the last move, the balancer will stop waiting for a block to be moved.
dfs.block.invalidate.limit	1000	The maximum number of invalid blocks that can be in the system. This property is only used when the "dfs.namenode.invalidate.block" property is set to true.
dfs.block.misreplication.processing.limit	10000	Maximum number of blocks that can be in the system that are being processed by the balancer.
dfs.block.replicator.classname	org.apache.hadoop.hdfs.server.blockmanagement.BlockPlacementPolicyDefault	Class representing block placement policies currently in use. The class org.apache.hadoop.hdfs.server.blockmanagement.BlockPlacementPolicyDefault will load the JSON file defined in the namenode configuration file "dfsadmin -refreshNodes" on the namenode.

		BlockPlacementPolicyWith domain policy. The details of corresponding policy classes corresponding class is org.apache.hadoop.hdfs.serv
dfs.blockreport.incremental.intervalMsec	0	If set to a positive integer, th reports from the Datanode to
dfs.checksum.type	CRC32C	Checksum type
dfs.client.block.write.locateFollowingBlock.retries	5	Number of retries to use wh
dfs.client.failover.proxy.provider		The prefix (plus a required r proxy provider for the host. "Configuration Details" sect
dfs.client.key.provider.cache.expiry	864000000	DFS client security key cach
dfs.client.max.block.acquire.failures	3	Maximum failures allowed v
dfs.client.read.prefetch.size		The number of bytes for the operation. Defaults to 10 * \$
dfs.client.read.short.circuit.replica.stale.threshold.ms	1800000	Threshold in milliseconds fc
dfs.client.read.shortcircuit.buffer.size	1048576	Buffer size in bytes for short
dfs.client.replica.accessor.builder.classes		Comma-separated classes fo will use external BlockRead
dfs.client.retry.interval-ms.get-last-block-length	4000	Retry interval in millisecond datanodes.
dfs.client.retry.max.attempts	10	Max retry attempts for DFSO
dfs.client.retry.policy.enabled	false	If true, turns on DFSCClient r
dfs.client.retry.policy.spec	10000,6,60000,10	Set to pairs of timeouts and :
dfs.client.retry.times.get-last-block-length	3	Number of retries for calls to
dfs.client.retry.window.base	3000	Base time window in ms for extended linearly (e.g. 3000 9000 ms for third retry, etc.)
dfs.client.socket-timeout	60000	Default timeout value in mil
dfs.client.socketcache.capacity	16	Socket cache capacity (in en
dfs.client.socketcache.expiryMsec	3000	Socket cache expiration for
dfs.client.test.drop.namenode.response.number	0	The number of Namenode re testing the NN retry cache.
dfs.client.hedged.read.threadpool.size	0	Support 'hedged' reads in DI number. The threadpool size 'hedged', concurrent reads in
dfs.client.hedged.read.threshold.millis	500	Configure 'hedged' reads in starting up a 'hedged' read.
dfs.client.use.legacy.blockreader	false	If true, use the RemoteBlock RemoteBlockReader2 class.
dfs.client.write.byte-array-manager.count-limit	2048	The maximum number of ar
dfs.client.write.byte-array-manager.count-reset-time-period-ms	10000	The time period in milliseco zero if there is no increment
dfs.client.write.byte-array-manager.count-threshold	128	The count threshold for each allocation count exceeds the managed until the allocation
dfs.client.write.byte-array-manager.enabled	false	If true, enables byte array m
dfs.client.write.max-packets-in-flight	80	The maximum number of D
dfs.content-summary.limit	5000	The maximum content sumr number means no limit (i.e.
dfs.content-summary.sleep-microsec	500	The length of time in micros in content summary comput
dfs.data.transfer.client.tcpnodelay	true	If true, set TCP_NODELAY
dfs.datanode.balance.max.concurrent.moves	50	Maximum number of thread reconfigurable via the "dfsac
dfs.datanode.fsdataset.factory		The class name for the unde org.apache.hadoop.hdfs.serv
dfs.datanode.fsdataset.volume.choosing.policy		The class name of the policy org.apache.hadoop.hdfs.serv you would like to take into a "org.apache.hadoop.hdfs.ser
dfs.datanode.hostname		Optional. The hostname for different for each machine. I
dfs.datanode.lazywriter.interval.sec	60	Interval in seconds for Datai
dfs.datanode.network.counts.cache.max.size	2147483647	The maximum number of en contain.
dfs.datanode.oob.timeout-ms	1500,0,0,0	Timeout value when sending OOB_RESTART, OOB_RE

		respectively. Currently, only
dfs.datanode.parallel.volumes.load.threads.num		Maximum number of thread number of storage directorie
dfs.datanode.ram.disk.replica.tracker		Name of the class implemen org.apache.hadoop.hdfs.serv
dfs.datanode.restart.replica.expiration	50	During shutdown for restart,
dfs.datanode.socket.reuse.keepalive	4000	The window of time in ms b second request occurs within
dfs.datanode.socket.write.timeout	480000	Timeout in ms for clients so
dfs.datanode.sync.behind.writes.in.background	false	If set to true, then sync_file_ only valid when the property
dfs.datanode.transferTo.allowed	true	If false, break block transfer chunks.
dfs.ha.fencing.methods		A list of scripts or Java class failover. See the HDFS High configuration.
dfs.ha.standby.checkpoints	true	If true, a NameNode in Stan saves it to its local storage a
dfs.ha.zkfc.port	8019	The port number that the zoo
dfs.http.port		The http port for used for H
dfs.https.port		The https port for used for H
dfs.journalnode.edits.dir	/tmp/hadoop/dfs/journalnode/	The directory where the jour
dfs.journalnode.kerberos.internal.spnego.principal		Kerberos SPNEGO principa
dfs.journalnode.kerberos.principal		Kerberos principal name for
dfs.journalnode.keytab.file		Kerberos keytab file for the
dfs.ls.limit	1000	Limit the number of files pri DFS_LIST_LIMIT_DEFAU
dfs.mover.movedWinWidth	5400000	The minimum time interval, again.
dfs.mover.moverThreads	1000	Configure the balancer's mo
dfs.mover.retry.max.attempts	10	The maximum number of re
dfs.mover.max-no-move-interval	60000	If this specified amount of ti DataNode, on more effort w Mover iteration.
dfs.namenode.audit.log.async	false	If true, enables asynchronou
dfs.namenode.audit.log.token.tracking.id	false	If true, adds a tracking ID fo
dfs.namenode.available-space-block-placement-policy.balanced-space-preference-fraction	0.6	Only used when the dfs.bloc org.apache.hadoop.hdfs.serv Special value between 0 and Datanodes with less disk spe
dfs.namenode.backup.dnrpc-address		Service RPC address for the
dfs.namenode.delegation.token.always-use	false	For testing. Setting to true al is disabled.
dfs.namenode.edits.asynclogging	true	If set to true, enables asynch Namenode uses the tradition
dfs.namenode.edits.dir.minimum	1	dfs.namenode.edits.dir inclu dfs.namenode.edits.dir.requi directories must be greater th directories falls below dfs.na property defaults to 1.
dfs.namenode.edits.journal-plugin		When FSEditLog is creating encounters a URI with a sch class from "dfs.namenode.ec JournalManager and have a
dfs.namenode.file.close.num-committed-allowed	0	Normally a file can only be to a positive integer N, a file complete.
dfs.namenode.inode.attributes.provider.class		Name of class to use for del
dfs.namenode.inode.attributes.provider.bypass.users		A list of user principals (in s external attributes provider v stored in HDFS instead of th be returned when requested.
dfs.namenode.max-num-blocks-to-log	1000	Puts a limit on the number o report.
dfs.namenode.max.op.size	52428800	Maximum opcode size in by
dfs.namenode.name.cache.threshold	10	Frequently accessed files tha FSDirectory nameCache.
dfs.namenode.replication.max-streams	2	Hard limit for the number of
dfs.namenode.replication.max-streams-hard-limit	4	Hard limit for all replication

dfs.namenode.replication.pending.timeout-sec	-1	Timeout in seconds for block replication pending.
dfs.namenode.stale.datanode.minimum.interval	3	Minimum number of missed heartbeat from a DataNode. The actual interval (dfs.namenode.stale.datanode.minimum.interval) value is greater than the projected value above is used.
dfs.namenode.snapshot.capture.openfiles	false	If true, snapshots taken will capture open files. Even after the capture, the snapshot will be valid leases. Note: The file length cannot be greater than the previous point-in-time version. Note: The file length cannot be greater than the previous point-in-time version. NameNode at the time of snapshot. In order to capture the snapshot, the file length cannot be greater than the previous point-in-time version. SyncFlag.UPDATE_LENGTH.
dfs.namenode.snapshot.skip.capture.accesstime-only-change	false	If accessTime of a file/directory is greater than the previous point-in-time version, the changed access time is not captured. If there is other modification together with the modification, the access time is captured.
dfs.pipeline.ecn	false	If true, allows ECN (explicit congestion notification) for pipeline.
dfs.qjournal.accept-recovery.timeout.ms	120000	Quorum timeout in milliseconds for specific segment.
dfs.qjournal.finalize-segment.timeout.ms	120000	Quorum timeout in milliseconds for specific segment.
dfs.qjournal.get-journal-state.timeout.ms	120000	Timeout in milliseconds when getting journal state.
dfs.qjournal.new-epoch.timeout.ms	120000	Timeout in milliseconds when creating a new epoch.
dfs.qjournal.prepare-recovery.timeout.ms	120000	Quorum timeout in milliseconds for specific segment.
dfs.qjournal.queued-edits.limit.mb	10	Queue size in MB for quorum journal.
dfs.qjournal.select-input-streams.timeout.ms	20000	Timeout in milliseconds for selecting input streams.
dfs.qjournal.start-segment.timeout.ms	20000	Quorum timeout in milliseconds for specific segment.
dfs.qjournal.write-txns.timeout.ms	20000	Write timeout in milliseconds for specific segment.
dfs.quota.by.storage.type.enabled	true	If true, enables quotas based on storage type.
dfs.secondary.namenode.kerberos.principal		Kerberos principal name for secondary namenode.
dfs.secondary.namenode.keytab.file		Kerberos keytab file for the secondary namenode.
dfs.support.append	true	Enables append support on the secondary namenode.
dfs.web.authentication.filter	org.apache.hadoop.hdfs.web.AuthFilter	Authentication filter class used for web authentication.
dfs.web.authentication.simple.anonymous.allowed		If true, allow anonymous user authentication.
dfs.web.ugi		dfs.web.ugi is deprecated. Use dfs.web.authentication.filter instead.
dfs.webhdfs.netty.high.watermark	65535	High watermark configuration for Netty.
dfs.webhdfs.netty.low.watermark	32768	Low watermark configuration for Netty.
dfs.webhdfs.oauth2.access.token.provider		Access token provider class for OAuth2.
dfs.webhdfs.oauth2.client.id		Client id used to obtain access token.
dfs.webhdfs.oauth2.enabled	false	If true, enables OAuth2 in WebHDFS.
dfs.webhdfs.oauth2.refresh.url		URL against which to post for token refresh.
ssl.server.keystore.keypassword		Keystore key password for HTTPS.
ssl.server.keystore.location		Keystore location for HTTPS.
ssl.server.keystore.password		Keystore password for HTTPS.
dfs.balancer.keytab.enabled	false	Set to true to enable login using keytab.
dfs.balancer.address	0.0.0.0:0	The hostname used for a keytab with dfs.balancer.keytab.enabled.
dfs.balancer.keytab.file		The keytab file used by the balancer. If configured with dfs.balancer.keytab.enabled, the balancer will use this keytab file.
dfs.balancer.kerberos.principal		The Balancer principal. This Balancer will substitute _HOST placeholder allows to use Keytab based login can be enabled.
ssl.server.truststore.location		Truststore location for HTTPS.
ssl.server.truststore.password		Truststore password for HTTPS.
dfs.lock.suppress.warning.interval	10s	Instrumentation reporting log interval.
dfs.webhdfs.use.ipc.callq	true	Enables routing of webhdfs requests through the IPC call queue.
httpfs.buffer.size	4096	The size buffer to be used when reading from the file system.
dfs.datanode.disk.check.min.gap	15m	The minimum gap between disk checks. The setting supports multiple time units: s, m, h, d. If the suffix is specified then milliseconds.



dfs.datanode.disk.check.timeout	10m	Maximum allowed time for does not complete within thi supports multiple time unit s specified then milliseconds i
dfs.use.dfs.network.topology	true	Enables DFSNetworkTopolo
dfs.qjm.operations.timeout	60s	Common key to set timeout supports multiple time unit s specified then milliseconds i
dfs.reformat.disabled	false	Disable reformat of NameN exist then attempt to format